

UrbanPulse — Combined Report (Tasks A, B, C)

DSE ZG556 / CC ZG556 — Stream Processing & Analytics

Situated Learning Assignment · Domain 3: Smart Cities & Urban Infrastructure · 75 marks

Real-Time Urban Operations Intelligence Platform for MetroConnect. Every result shown was produced by running the accompanying code against the Docker Compose stack.

Task A — Architecture Design

Lambda vs Kappa for a Dual-Reporting City Platform

UrbanPulse — Real-Time Urban Operations Intelligence Platform Course: DSE ZG556 / CC ZG556 — Stream Processing & Analytics · Marks: 20 (M1 + M2)

A.0 Executive Summary


MetroConnect requires **two things at once** that pull architecture in opposite directions:

1. **Real-time operational response** — AQI emergency alerts in < 2 minutes, adaptive signal control within 90 seconds, live bus ETAs refreshed in < 60 seconds.
2. **Authoritative batch reporting** — weekly/monthly ward-level reports that are submitted to elected councillors and the state government, and therefore must be **reproducible and auditable**.

This document presents the UrbanPulse reference architecture, a formal Lambda-vs-Kappa evaluation grounded in MetroConnect's specific streams and mandates, and a 16-item government deployment readiness checklist. **The recommended choice is a *lean Lambda* architecture** — the reasoning is in §A.3.

A.1 UrbanPulse Architecture Diagram

All four data sources → a 3-broker Kafka ingestion backbone → **both** a real-time (speed) processing layer **and** a batch processing layer → a polyglot storage layer with a purpose-chosen database per data class → a serving layer of dashboards, advisory APIs and the signal-control write-back interface.



A.1.1 Storage technology choices (with justification)

The open-source mandate and data-sovereignty rule (data must not leave city servers) constrain every choice to a **self-hostable, on-premise, open-source** engine.

Data class (requirement)	Chosen technology	Why this technology
Time-series sensor data — raw AQI, signal metrics, meter readings; high write rate, time-ordered, needs downsampling	TimescaleDB (PostgreSQL + hypertables)	Native time-partitioned hypertables, automatic compression, <i>continuous aggregates</i> for downsampling, and plain SQL — so city IT staff already fluent in Postgres can operate it. Open-source, runs entirely on city servers.
Geospatial bus positions — live "nearest bus / buses on a corridor", sub-second, high churn	Redis GEO (hot layer) + PostGIS (analytical)	Redis GEOADD / GEOSearch gives in-memory sub-second radius queries for the live map & bunching lookups, with TTL so stale positions self-expire. PostGIS handles richer route/corridor spatial analytics on positions persisted for history.
Historical AQI records — 90-day+ retention, pollution-trend analytics over months	ClickHouse (columnar) with Parquet-on-MinIO as cold data lake	ClickHouse gives fast columnar rollups (min/max/percentile AQI per zone over months); MinIO+Parquet is the cheap, immutable, S3-API-compatible cold store that never leaves the city DC.
Councillor report aggregates — pre-computed ward/zone rollups feeding official dashboards	ClickHouse serving mart → Apache Superset	Aggregates are small, read-heavy, and dashboard-facing; ClickHouse serves them in milliseconds and Superset gives councillors/ward officers no-code, accessible dashboards.

Why both a speed layer and a batch layer are shown: the speed layer (Flink + Spark Structured Streaming) satisfies the sub-2-minute operational SLAs; the batch layer (Spark over the immutable MinIO/Parquet lake, scheduled by Airflow) produces the *authoritative, reproducible* figures that go into government submissions. This duality is the crux of the Lambda-vs-Kappa decision below.

A.2 Lambda vs Kappa — Evaluation Matrix (UrbanPulse-grounded)

Every cell is assessed against MetroConnect's actual streams, SLAs and mandates — not generic textbook properties.

Criterion	Lambda (speed layer + batch layer)	Kappa (single stream engine, replay for "batch")
Latency	Speed layer (Flink) hits the hard SLAs directly — AQI alert < 2 min, signal adaptation < 90 s; the batch layer adds hours-scale latency for councillor reports, which is fine because those are weekly/monthly. Both latency classes met, via two paths.	Single Flink pipeline also hits < 2 min / 90 s for real-time. "Batch" councillor figures come from long-window aggregation or Kafka replay — acceptable since reporting latency is relaxed. Latency parity with Lambda for the real-time path.
Fault Tolerance	Two failure domains to keep alive. Upside: the batch layer recomputes from the immutable Parquet lake, so a corrupted live result is self-healing on the next run; the speed layer can tolerate approximate results between batch corrections.	One pipeline with Flink checkpoints + Kafka replay → fewer components, fewer failure domains. Risk: a single bad deploy simultaneously degrades <i>both</i> live alerts and the reporting figures — no independent batch safety net.
Operational Complexity	Highest cost of Lambda: two codebases (Flink speed + Spark batch) compute overlapping metrics (e.g., ward AQI). Logic can diverge — and a councillor report disagreeing with the live dashboard is politically damaging. Heavier for a budget-constrained municipal team.	One engine, one codebase (Flink) → far lower cognitive & staffing load, attractive for a small city team. But demands rigorous Kafka retention/replay discipline and schema governance to make replays trustworthy.
Reprocessing Capability	Reprocessing = re-run a Spark batch job over the Parquet lake — mature and cheap for historical corrections (e.g., recomputing 6 months of AQI after a sensor recalibration). Scans columnar Parquet, not the event log.	Reprocessing = replay the Kafka log through a new job version. Our retention (AQI 90 d, meters 365 d) supports it, but replaying 365 days of meter data through Flink is far heavier than a Spark scan of pre-compacted Parquet.
Cost	Two clusters (stream + batch) → more compute; but batch runs off-peak and the cold lake (MinIO/Parquet) is cheap bulk storage. Predictable, but a larger footprint.	One cluster is cheaper to run, but Kappa must retain long Kafka history (365 d meters) on broker storage to stay replayable → Kafka storage cost grows; needs tiered storage to contain it.
Compliance with Government Reporting Mandate	Decisive strength. The batch layer produces official figures by a <i>fixed, versioned, re-runnable</i> computation over immutable stored data — exactly the determinism, provenance and auditability that state-government submissions and auditors require.	Official figures derive from stream replay; reproducibility hinges on exactly-once + retention + deterministic operators. Auditors may distrust "streaming-generated" statutory numbers unless replay determinism is formally proven and frozen.

A.3 Architecture Choice & Justification

Chosen: a *lean Lambda* architecture.

- **Speed layer:** Apache Flink for incident detection (AQI/gridlock/bunching) and Spark Structured Streaming for the 15-minute ward energy and rolling-AQI advisory windows.
- **Batch layer:** Spark batch jobs over an immutable MinIO/Parquet data lake, scheduled by Airflow, producing the authoritative weekly/monthly councillor and state-government rollups.
- **"Lean" mitigation of Lambda's chief weakness (duplicate logic):** the batch and speed layers share metric definitions through common Spark SQL / UDFs, and the batch set is a small, stable list

of ward/zone rollups — so divergence risk and maintenance burden stay low.

Why Lambda over Kappa here — the deciding factor is the *government reporting mandate*. For most commercial streaming systems Kappa wins on simplicity, and if UrbanPulse were *only* an operations platform I would choose Kappa. But MetroConnect's reports are **statutory submissions to elected councillors and the state government**. Those figures must be reproducible on demand, auditable, and defensible months later — properties that a batch recomputation over an immutable data lake provides cleanly and that stream-replay provides only with significant additional proof-of-determinism engineering. The real-time SLAs (< 2 min AQI, 90 s signals) are met either way, so they do not break the tie; the audit/compliance requirement does. Hence: **Lambda**.

A.4 Architecture Readiness Checklist — Government Smart-City Deployment

16 items (minimum 12 required). Each is a concrete, checkable go-live gate. Grouped by the four mandated concern areas.



Data Sovereignty (data must not leave city servers)

1. **On-premise only** — all Kafka brokers, processing engines, and databases run inside the city data centre (or a state-government cloud region physically within the state); no managed SaaS that egresses data outside city jurisdiction.
2. **No external egress** — network egress firewall rules deny all outbound connections from the data plane except explicitly whitelisted intra-DC endpoints; verified by egress audit.
3. **Data residency attestation** — every storage volume (TimescaleDB, ClickHouse, MinIO, Redis) is confirmed to reside on city-owned/leased hardware within the municipal boundary, with a signed residency record.
4. **Encryption in transit & at rest** — TLS on all Kafka listeners and DB connections; at-rest encryption on MinIO and DB volumes, with keys held in a city-controlled KMS/Vault.



Open-Source Mandate

5. **100% OSS stack, licence-cleared** — Kafka, Flink, Spark, TimescaleDB, ClickHouse, Redis, PostGIS, MinIO, Superset/Grafana are all OSI-approved licences; a licence inventory (SBOM) is reviewed for copyleft/attribution obligations — no proprietary lock-in.
6. **No vendor-locked formats** — data is stored in open formats (Parquet, JSON, SQL) so the city can migrate engines without re-ingesting from source.
7. **Reproducible builds** — all services pinned to specific versions and built from published images/Dockerfiles committed to the city's git repo (see this project's `docker-compose.yml`), so the deployment is auditable and rebuildable.



Disaster Recovery (RPO < 15 min, RTO < 30 min)

8. **RPO < 15 min** — Kafka replication factor = 3 with `min.insync.replicas = 2`, plus TimescaleDB/ClickHouse WAL/replication and MinIO versioning snapshotted at ≤ 10-minute intervals, so at most < 15 minutes of data can be lost.

9. **RTO < 30 min** — a documented, *drill-tested* failover runbook (broker/quorum recovery, engine restart from last checkpoint, DB promote-replica) that has been demonstrated to restore service in under 30 minutes.
10. **Flink/Spark checkpointing** — exactly-once checkpoints persisted to replicated storage so stream jobs resume from the last checkpoint (bounded reprocessing) rather than from zero after a failure.
11. **Cross-rack / cross-site redundancy** — brokers and storage spread across ≥ 2 racks (ideally 2 sites within the city) so a single rack/PSU/switch failure cannot take the quorum below majority.
12. **Backup restore verified** — periodic restore rehearsals prove backups are actually recoverable (a backup that has never been restored is not a backup).

Accessibility for Non-Technical Ward Officers

13. **No-code dashboards** — ward officers use Apache Superset/Grafana dashboards with plain-language labels (e.g., "Air Quality: Unhealthy in Old City") — no SQL or query language required to read the operational picture.
14. **Localisation & readability** — dashboards available in the regional language plus English, with colour-blind-safe palettes and mobile/tablet-responsive layouts for field use.
15. **Role-based access & simple alerts** — ward officers see only their ward/zone by default; critical advisories (AQI breach, gridlock) are pushed as plain SMS/app notifications, not buried in a technical console.
16. **WCAG-aligned accessibility & training** — dashboards meet basic WCAG 2.1 AA contrast/keyboard requirements, and every ward officer completes a short onboarding so the system is usable by non-technical staff on day one.

End of Task A.

Task B — Apache Kafka: Multi-Source Urban Data Ingestion

UrbanPulse · DSE ZG556 Stream Processing & Analytics · Marks: 20 (M2)

All results shown below were produced by running the code in this repository against the 3-broker KRaft cluster in `docker-compose.yml`. Commands to reproduce each result are given inline. Source lives in `src/task_b_kafka/`.

B.0 Component Map

#	Requirement	Module	Evidence (§)
1	3-broker cluster + retention policies	<code>docker-compose.yml</code> , <code>create_topics.py</code>	B.1
2	Producers — at-least-once, route-key ordering, null handling	<code>producer_bus_gps.py</code> , <code>producer_air_quality.py</code>	B.2
3	Priority consumer architecture (HIGH zero-lag)	<code>priority_consumers.py</code>	B.3
4	Kafka Streams route enrichment (KTable join)	<code>load_route_schedule.py</code> + <code>task_b_kafka_streams/</code>	B.4
5	Dead-letter queue + 5-min report	<code>dlq_router.py</code> , <code>dlq_report.py</code>	B.5

B.1 Cluster Setup & Retention Policies (4 marks)

Cluster. Three brokers in **KRaft mode** (no ZooKeeper), each acting as both broker and controller, forming a 3-voter metadata quorum. Replication factor is **3** with `min.insync.replicas = 2`, so the cluster survives the loss of one broker with no data loss when producers use `acks=all`.

```
$ docker exec urbanpulse-broker1 kafka-metadata-quorum --bootstrap-server broker1:19092 describe --status
ClusterId: MKU30EVBNTcwNTJENDM2Qk   LeaderId: 2   CurrentVoters: [1,2,3]
MaxFollowerLag: 0
```

Topics, partitions & retention (`python -m src.task_b_kafka.create_topics`):

Topic	Partitions	Retention	Justification
urbanpulse.bus_gps	12	24 h	Highest rate (~2,400 ev/s); keyed by <code>route_id</code> . 12 partitions ≈ 200 ev/s each and up to 12 parallel consumers. 24 h = replay bus positions during accident/incident investigations.
urbanpulse.traffic_signals	6	7 d	Keyed by <code>junction_id</code> . 6 is divisible by the STANDARD consumer-group size (3) for even assignment and read whole by the 1-consumer HIGH group. Signal history ages out fast → 7 d.
urbanpulse.air_quality	3	90 d	Low rate (~60 ev/s), keyed by <code>zone</code> . 3 partitions give modest parallelism without over-partitioning. 90 d = seasonal pollution-trend analysis.
urbanpulse.smart_meters	12	365 d	High rate (~1,100 ev/s), keyed by <code>ward_id</code> (12 wards) → one partition per ward. 365 d = annual regulatory energy audit.

Three distinct retention values (24 h / 90 d / 365 d), each tied to a concrete operational or regulatory driver, exactly as required. Verified live:

```
$ python -m src.task_b_kafka.create_topics --describe
urbanpulse.bus_gps      partitions=12  retention.ms=86400000    (~1d)   min.isr=2
urbanpulse.air_quality  partitions=3   retention.ms=7776000000  (~90d)  min.isr=2
urbanpulse.smart_meters partitions=12   retention.ms=31536000000 (~365d) min.isr=2
```

B.2 Producers (5 marks)

Bus GPS — ordered per route (`producer_bus_gps.py`)

Ordering of positions **per route** is guaranteed by two mechanisms together:

1. **key = route_id** → the default partitioner sends every event for a route to the *same* partition, and Kafka guarantees order within a partition.
2. **enable.idempotence = True + acks=all** → per-partition order is preserved even across retries and with up to 5 in-flight batches, and producer retries are de-duplicated.

```
$ python -m src.task_b_kafka.producer_bus_gps --count 2000
[bus_gps] producing to urbanpulse.bus_gps (key=route_id, idempotent, acks=all)
[bus_gps] sent=2000 delivered=2000 failed=0
```


Air Quality — at-least-once + retry + null handling (producer_air_quality.py)

- **At-least-once:** `acks=all` + client `retries=5`; `enable.idempotence=False` so a retried send may duplicate — acceptable under at-least-once (a missed AQI breach is worse than a duplicate).
- **Explicit retry logic** (sensors "occasionally timeout"): a delivery that fails after the client's own retries is re-queued by the application up to `MAX_APP_RETRIES=3` times, tracked via a message header `attempt`.
- **Null-AQI handling:** the simulator injects the mandated ~5% null-AQI sensor faults. The producer logs and counts each one and forwards it (never crashes), so the downstream DLQ can classify it. Graceful = observed, counted, non-fatal.

```
$ python -m src.task_b_kafka.producer_air_quality --count 2000
[air_quality] ▲ null AQI from AQ-Z01-00 (zone=Z01-CentralBusiness) – forwarding for DLQ
[total nulls=25]
[air_quality] sent=2000 delivered=2000 app-retries=0 perm-failed=0
[air_quality] null-AQI handled gracefully: 94 (4.7%) ← ≈ the injected 5%
```

B.3 Priority Consumer Architecture (6 marks)

Two consumer groups read the **same** `traffic_signals` topic with independent committed offsets:

Group	Members	Partitions each	Role
urbanpulse-signals-high (HIGH_PRIORITY)	1	all 6	real-time signal control — must stay ~0 lag
urbanpulse-signals-standard (STANDARD_PRIORITY)	3	2	analytics dashboard — lag-tolerant

Because the groups have **independent offsets**, a deliberate slowdown in the STANDARD group (150 ms/msg) cannot slow HIGH. The self-contained demo produces signals at 300 ev/s and prints live lag for both groups:

```
$ python -m src.task_b_kafka.priority_consumers --role demo --duration 30 --slow-ms 150 -
-prod-rate 300
```

time	HIGH lag	HIGH done	STD lag	STD done	(topic has 6 partitions)
0s	0	167	156	15	
9s	7	2,199	2,007	195	
18s	8	4,235	3,864	373	
27s	13	6,244	5,695	552	

Result: HIGH_PRIORITY lag stays in single digits (near-zero) and keeps draining all partitions, while STANDARD_PRIORITY lag climbs past **5,600** because its 3 slow consumers process only ~20 ev/s combined against 300 ev/s of production. The critical signal-control path is fully isolated from analytics back-pressure — exactly the guarantee the city needs for 90-second adaptive signal control.

B.4 Kafka Streams Route Enrichment (3 marks)

This is implemented as an **actual Apache Kafka Streams** application in Java (`src/task_b_kafka_streams/RouteEnrichmentApp.java`), rather than a Python look-up loop. The supplied `route_schedule.csv` is first loaded by `load_route_schedule.py` into `urbanpulse.route_schedule`, a **compacted topic** keyed by `route_id`. Kafka Streams materialises that topic as a local KTable and performs this left join:

```
bus_gps (KStream, keyed route_id) LEFT JOIN route_schedule (compacted KTable)
  → urbanpulse.bus_enriched (+ route_name, terminal, scheduled_arrival_time)
```

The service uses `exactly_once_v2` processing. A GPS event with an unknown `route_id` is emitted with `join_status = NO_ROUTE_MATCH` rather than dropped, making the exception visible for data-quality follow-up.

```
$ docker compose exec app python -m src.task_b_kafka.load_route_schedule
[route-schedule] loaded 20 rows into urbanpulse.route_schedule (key=route_id,
cleanup.policy=compact)
$ docker compose --profile streams up -d --build route-enrichment
[kafka-streams] KStream(bus_gps) LEFT JOIN KTable(route_schedule) →
urbanpulse.bus_enriched
```

Sample enriched record on `urbanpulse.bus_enriched`:

```
{ "bus_id": "BUS-R002-00", "route_id": "R002", "lat": 12.869167, "lon": 77.541086,
  "speed_kmh": 42.4, "occupancy_pct": 26, "timestamp": 1783647584927,
  "route_name": "Shivajinagar - Whitefield", "terminal": "Whitefield TTMC",
  "scheduled_arrival_time": "08:20", "join_status": "MATCHED" }
```

This enriched stream (position + schedule) is the foundation for the real-time ETA service.

B.5 Dead-Letter Queue (2 marks)

`dlq_router.py` consumes all four raw streams, validates each event against the per-stream rules in `src/common/schemas.py`, and routes failures to `urbanpulse.dlq` with an `error_reason`, the source topic/partition/offset, and the original payload (for replay). **Well over 3 validation rules** are enforced — e.g. `NULL_AQI`, `AQI_OUT_OF_RANGE`, `IMPOSSIBLE_GPS`, `IMPOSSIBLE_SPEED`, `NULL_WAIT`, `VOLTAGE_OUT_OF_RANGE`, `POWER_FACTOR_OUT_OF_RANGE`, `NULL_KWH`, `MALFORMED_JSON`.

```
$ python -m src.task_b_kafka.dlq_router --from-beginning --duration 25
[dlq] FINAL: seen=14,012 rejected=244
```

5-minute DLQ report (`dlq_report.py`) — error-type distribution:

error_reason	count	share
NULL_AQI	107	43.9%
NULL_WAIT	84	34.4%
IMPOSSIBLE_GPS	50	20.5%
VOLTAGE_OUT_OF_RANGE	3	1.2%

stream	rejected	top reasons
air_quality	107	NULL_AQI:107
traffic_signals	84	NULL_WAIT:84
bus_gps	50	IMPOSSIBLE_GPS:50
smart_meters	3	VOLTAGE_OUT_OF_RANGE:3

The rejection rate (~1.7% overall) matches the anomaly rates the simulators inject (5% null AQI, ~1.5% impossible GPS, ~1% null wait, ~0.5% bad voltage), confirming the validation and routing are correct end-to-end.

B.6 How to Run (all of Task B)

```

docker compose up -d                                # 3 brokers + Kafka-UI + app
docker compose exec app python -m src.task_b_kafka.create_topics
# producers
docker compose exec app python -m src.task_b_kafka.producer_bus_gps --rate 120 --duration
60
docker compose exec app python -m src.task_b_kafka.producer_air_quality --rate 20 --
duration 60
# priority consumers (self-contained demo)
docker compose exec app python -m src.task_b_kafka.priority_consumers --role demo
# enrichment join
docker compose exec app python -m src.task_b_kafka.load_route_schedule
docker compose --profile streams up -d --build route-enrichment
# DLQ
docker compose exec app python -m src.task_b_kafka.dlq_router --from-beginning --duration
60
docker compose exec app python -m src.task_b_kafka.dlq_report

```

Kafka-UI at <http://localhost:8080> shows all topics, partitions, consumer-group lag, and messages live — useful for the video walkthrough.

End of Task B.

Task C — Flink Real-Time Incident Detection & Spark Urban Analytics Engine

UrbanPulse · DSE ZG556 Stream Processing & Analytics · Marks: 35 (M3)

Code in `src/task_c_flink_spark/`. The Flink and Spark runtimes are provided as Docker images (`infra/flink.Dockerfile` , `infra/spark.Dockerfile`) wired into `docker-compose.yml` under the `flink` and `spark` profiles. Verified outputs are shown inline; reproduce with the commands in §C.5.

C.0 Component Map

#	Requirement	Module	Marks
I	Flink — 3 incident patterns, keyed state + event-time watermarks	<code>flink_incident_detection.py</code>	10
II.a	Spark — ward energy 15-min tumbling window, 45-min watermark, Kafka+Parquet	<code>spark_ward_energy.py</code>	20
II.b	Spark — 10-min rolling-avg AQI SQL, <code>zone_profile</code> join, Update mode	<code>spark_health_advisory.py</code>	(within II)
III	Flink vs Spark — use-case mapping across 4 dimensions	this doc §C.4	5

C.1 Part I — Apache Flink Incident Detection (10 marks)

A single Flink DataStream job (`flink_incident_detection.py`) runs three detectors, each keyed and using event-time watermarks (5 s bounded out-of-orderness on the events' `timestamp` field). Checkpointing every 30 s makes the keyed state recoverable. All alerts are unioned into one stream and written to `urbanpulse.incidents`.

Pattern	Key	State	Trigger
(a) AQI Emergency	sensor_id	ValueState<Long> last-alert time	aqi > 300 ; 2-min event-time cooldown suppresses duplicates
(b) Traffic Gridlock	junction_id	ValueState<Int> consecutive-cycle counter	avg_wait_sec > 180 for 3 consecutive cycles → alert(junction_id, zone), then reset
(c) Bus Bunching	route_id	ValueState<String> containing positions, pair-first-seen times and alert flags	two buses < 200 m apart for > 5 min; a checkpointed event-time timer emits alert(bus_a, bus_b) at the threshold

Why keyed state + event time here: each detector's logic is intrinsically per-entity and temporal.

Gridlock's "3 consecutive cycles" is a per-junction counter; bunching's "within 200 m for 5 minutes" is a per-route, per-pair event-time timer registered through `ctx.timer_service()`. The timer is checkpointed with the state and fires once the watermark passes the threshold, so an alert does not depend on another event from that same pair arriving after five minutes. Event-time watermarks keep the calculation correct despite slightly out-of-order readings.

Events with null critical fields are dropped at parse time (they are handled by the Task B DLQ), so the detectors operate on clean data.

Verified output — all three detectors fire and alerts are written to `urbanpulse.incidents` (bunching window shortened to 20 s for the demo via `URBANPULSE_BUNCHING_SECONDS`; the code default is the full 5 minutes):

```
=== incident types (one 70s demo run) ===
11 AQI_EMERGENCY
44 TRAFFIC_GRIDLOCK
1  BUS_BUNCHING

=== one sample of each (JSON on urbanpulse.incidents) ===
{"incident_type":"AQI_EMERGENCY","severity":"HAZARDOUS","sensor_id":"AQ-Z02-02",
 "zone":"Z02-IndustrialEast","aqi":395.8,"detail":"AQI 396 > 300 at AQ-Z02-02 (Z02-IndustrialEast)"}
{"incident_type":"TRAFFIC_GRIDLOCK","junction_id":"JN-Z05-01","zone":"Z05-OldCity",
 "avg_wait_sec":220.8,"consecutive_cycles":3,"detail":"JN-Z05-01 (Z05-OldCity) avg_wait 221s for 3 consecutive cycles"}
{"incident_type":"BUS_BUNCHING","route_id":"R001","bus_a":"BUS-R001-00","bus_b":"BUS-R001-01",
 "distance_m":44.3,"held_seconds":20.0,"detail":"BUS-R001-00 & BUS-R001-01 on R001 within 44m for >20s"}

$ kafka-get-offsets --topic urbanpulse.incidents → 264 incidents persisted
```

Keyed-state note (engineering detail). The bunching detector keeps its per-route state (positions, per-pair proximity timers, alerted flags) in a single `ValueState<String>` JSON blob keyed by `route_id`, with 5-minute position pruning to bound the state. This proved more robust in PyFlink than

a non-STRING-valued `MapState`. The gridlock and AQI detectors use `ValueState<Int>` / `ValueState<Long>` respectively.

C.2 Part II — Spark Structured Streaming: Ward Energy (20 marks)

`spark_ward_energy.py` consumes `urbanpulse.smart_meters` and computes, per `ward_id` per **15-minute tumbling window** with a **45-minute late-data watermark**:

```
parsed.withWatermark("event_time", "45 minutes")
    .groupBy(window("event_time", "15 minutes"), "ward_id")
    .agg(sum("kwh_reading")      .alias("total_kwh_consumed"),
         avg("power_factor")    .alias("avg_power_factor"),
         max("voltage")         .alias("peak_voltage"))
```

Dual sink (both required, two independent streaming queries on the same aggregated DataFrame):

1. **Kafka** → `urbanpulse.ward_energy_summary` (one JSON row per closed window).
2. **Parquet** → `data/parquet/ward_energy`, **partitioned by `ward_id` and `date`** for historical trend analysis.

Append output mode is used: a windowed aggregation with a watermark emits each window's final result once the watermark passes the window end — and append is also the only mode the Parquet file sink supports. The 45-minute watermark means a meter reading up to 45 minutes late is still counted in its correct window.

Verified output (demo window shortened to 30 s / 10 s watermark so windows close quickly; the graded defaults are 15 min / 45 min). All 12 wards emit once per closed window; injected sensor-fault voltages are filtered out so `peak_voltage` stays in the real 218–242 V band:

ward_id	window_start	window_end	total_kwh_consumed
avg_power_factor	peak_voltage		
-----	-----	-----	-----
----	-----	-----	-----
W01	2026-07-10 05:03:00	2026-07-10 05:03:30	17.709
0.9252	241.7		
W02	2026-07-10 05:03:00	2026-07-10 05:03:30	21.558
0.9228	241.9		
W03	2026-07-10 05:03:00	2026-07-10 05:03:30	19.775
0.9286	241.6		
...			
W12	2026-07-10 05:03:00	2026-07-10 05:03:30	13.772
0.9092	241.9		

The same rows are simultaneously published as JSON to `urbanpulse.ward_energy_summary` and appended to the Parquet lake below.

Partitioned Parquet written to disk (partitioned by `ward_id` then `date`):

```
data/parquet/ward_energy/
├─ ward_id=W01/date=2026-07-10/part-00000-....snappy.parquet
├─ ward_id=W02/date=2026-07-10/part-00002-....snappy.parquet
├─ ... (one directory per ward)           → 36 parquet files across 12 wards
└─ _spark_metadata/
```

C.2.1 Streaming SQL — AQI Health Advisories (part of Part II)

`spark_health_advisory.py` is a **Streaming SQL query** on `urbanpulse.air_quality`:

- (a) a **10-minute rolling average AQI per zone**, implemented as a sliding `window(event_time, '10 minutes', '1 minute')`,
- (b) **joined with the static `zone_profile` table** (zone name, population, number of schools) — a stream-static join,
- (c) **filtered for `rolling_avg_aqi > 150`** (Unhealthy) and written to `urbanpulse.health_advisories` in **Update** output mode.

The logic is expressed as an actual SQL string over temp views (see the module), so a ward officer receives "Unhealthy air over 620,000 people and 110 schools in Residential North", not a bare number.

Verified output (demo rolling window shortened to 1 minute and recomputed every 10 seconds) — enriched advisories on `urbanpulse.health_advisories`, Update mode re-emitting as each zone's rolling average evolves (note `readings` climbing 221 → 231 for the same window):

zone	zone_name	population	num_schools
rolling_avg_aqi	readings	advisory_level	
Z06-Airport	Airport & Aerotropolis	190000	27
170.8	221	UNHEALTHY	
Z03-ResidentialNorth	Residential North	620000	110
159.8	229	UNHEALTHY	
Z06-Airport	Airport & Aerotropolis	190000	27
170.1	223	UNHEALTHY	
Z03-ResidentialNorth	Residential North	620000	110
159.6	231	UNHEALTHY	

Only zones whose 10-minute rolling average exceeds 150 appear, each enriched with the population and school count a ward officer needs to gauge exposure.

C.4 Part III — Flink vs Spark for UrbanPulse (5 marks)

Mapping the two UrbanPulse use cases to the right engine, across the four required dimensions.

Dimension	Bus Bunching → Apache Flink	Ward Energy → Apache Spark
State size	Per-route keyed state holding each bus's latest position plus a timer per bus-pair , updated on every event. Many fine-grained, continuously-mutated keys — Flink's RocksDB keyed state with incremental checkpoints is built for exactly this.	A windowed aggregation keyed by ward (only 12 wards) over 15-min windows — small, bounded state. A classic, cheap windowed aggregation that Spark's structured state handles trivially.
Latency requirement	Bunching is a live operational incident — buses must be un-bunched quickly, so alerts need event-at-a-time, sub-second latency. Flink processes each event as it arrives.	Ward energy feeds councillor dashboards refreshed every 15 minutes — latency-relaxed. Spark's micro-batch (seconds) is far faster than needed; sub-second would be wasted effort.
Recovery time objective	Incident detection must resume fast after failure without losing the in-flight proximity timers. Flink's incremental RocksDB checkpoints + exactly-once restore large keyed state quickly → low RTO for a critical path.	A 15-min reporting job tolerates a slightly higher RTO; Spark checkpoints to reliable storage and simply reprocesses the current micro-batch on restart — perfectly adequate for non-critical analytics.
Operational complexity	The "within 200 m for > 5 min" rule is a per-pair, event-time state machine — natural in a Flink KeyedProcessFunction with timers , awkward in Spark (needs <code>flatMapGroupsWithState</code>). Flink is the simpler fit <i>for this logic</i> .	Windowed SUM/AVG/MAX + a dual sink including partitioned Parquet for a data lake is idiomatic, concise Spark SQL. Doing partitioned Parquet analytics output in Flink is more work. Spark is the simpler fit <i>for this logic</i> .

Conclusion. UrbanPulse deliberately uses **both** engines, each where it is strongest — which is exactly the *lean Lambda* architecture chosen in Task A: **Flink for low-latency, fine-grained, event-time incident detection** (AQI, gridlock, bunching) on the speed layer, and **Spark for windowed, SQL-friendly, data-lake-producing analytical aggregation** (ward energy, health advisories). Forcing either use case onto the other engine would increase either latency (Spark for bunching) or code complexity (Flink for partitioned Parquet reports).

C.5 How to Run (all of Task C)

```
# cluster + data
docker compose up -d
docker compose exec app python -m src.task_b_kafka.create_topics
docker compose exec -d app python -m src.simulators.run_simulator --all

# --- Flink incident detection (runs on the arm64-native Flink image) ---
docker compose --profile flink up -d
docker compose exec -e URBANPULSE_BUNCHING_SECONDS=45 flink-jobmanager \
    flink run -py /opt/job/src/task_c_flink_spark/flink_incident_detection.py
# ...or local mode: docker compose exec flink-jobmanager \
#     python /opt/job/src/task_c_flink_spark/flink_incident_detection.py

# --- Spark ward energy (dual sink) ---
docker compose --profile spark up -d
docker compose exec spark spark-submit \
```



```
--packages org.apache.spark:spark-sql-kafka-0-10_2.12:3.5.1 \
/opt/job/src/task_c_flink_spark/spark_ward_energy.py --duration 180
```

```
# --- Spark health advisories (streaming SQL) ---
```

```
docker compose exec spark spark-submit \
```

```
--packages org.apache.spark:spark-sql-kafka-0-10_2.12:3.5.1 \
```

```
/opt/job/src/task_c_flink_spark/spark_health_advisory.py \
```

```
--window "1 minute" --slide "10 seconds" --watermark "20 seconds" --duration 180
```

Watch results live in Kafka-UI (<http://localhost:8080>) on the `urbanpulse.incidents` , `urbanpulse.ward_energy_summary` , and `urbanpulse.health_advisories` topics, and the Flink dashboard at <http://localhost:8081> .

End of Task C.